

# SQL and Relational Algebra

Zaki Malik

September 02, 2008

# Basics of Relational Algebra

- Four types of operators:
  - Select/Show parts of a single relation: projection and selection.
  - Usual **set operations** (union, intersection, difference).
  - **Combine the tuples of two relations**, such as cartesian product and joins.
  - **Renaming**.

# Projection

- The projection operator produces from a relation R a new relation containing only some of R's columns.
- “Delete” (i.e. not show) attributes not in projection list.
- Duplicates eliminated
- To obtain a relation containing only the columns  $A_1, A_2, \dots, A_n$  of R

**RA:**  $\pi_{A_1, A_2, \dots, A_n} (R)$

**SQL:** `SELECT A1,A2, . . . An FROM R;`

# Selection

- The selection operator applied to a relation R produces a new relation with a subset of R's tuples.
- The tuples in the resulting relation satisfy some condition C that involves the attributes of R.
  - with duplicate removal

RA:  $\sigma_C(R)$

SQL: `SELECT *FROM R WHERE C;`

- The WHERE clause of a SQL command corresponds to  $\sigma( )$ .

# Union

- The union of two relations R and S is the set of tuples that are in **R or in S or in both**.
  - R and S must have **identical sets of attributes** and the types of the attributes must be the same.
  - The attributes of R and S must occur in the **same order**.
- What is the schema of the result ?

**RA:**  $R \cup S$

**SQL:** **(SELECT \* FROM R)**  
**UNION**  
**(SELECT \* FROM S);**

# Union

***S1***

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

***S2***

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

***S1***  $\cup$  ***S2***

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

# Intersection

- The intersection of two relations R and S is the set of tuples that are **in both R and S**.
- Same conditions hold on R and S as for the union operator.
  - R and S must have **identical sets of attributes** and the types of the attributes must be the same.
  - The attributes of R and S must occur in the **same order**.

**RA:**  $R \cap S$

**SQL:** **(SELECT \* FROM R)**  
**INTERSECT**  
**(SELECT \* FROM S);**

# Intersection

*S1*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

*S2*

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$S1 \cap S2$

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0



# Difference

- The difference of two relations R and S is the set of tuples that are **in R but not in S**.
- Same conditions hold on R and S as for the union operator.
  - R and S must have **identical sets of attributes** and the types of the attributes must be the same.
  - The attributes of R and S must occur in the **same order**.

**RA:**  $R - S$

**SQL:** **(SELECT \* FROM R)**  
EXCEPT  
**(SELECT \* FROM S);**

- $R - (R - S) = R \cap S$

# Difference

*S1*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

*S2*

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

*S1* − *S2*

sid	sname	rating	age
22	dustin	7	45.0

# Cartesian Product

- The Cartesian product (or cross-product or product) of two relations R and S is a the set of pairs that can be formed by pairing each tuple of R with each tuple of S.
  - The result is a relation whose schema is the schema for R followed by the schema for S.

RA:  $R \times S$

SQL: **SELECT \* FROM R , S ;**

# Cartesian Product

*S1*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

*R1*

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

*S1 x R1*

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

?

We **rename** attributes to avoid ambiguity or we **prefix attribute** with the name of the relation it belongs to.

# Theta-Join

- The theta-join of two relations R and S is the **set of tuples in the Cartesian product of R and S that satisfy some condition C.**

RA:  $R \underset{C}{\bowtie} S$

SQL: **SELECT \***  
**FROM R , S**  
**WHERE C;**

- $R \underset{C}{\bowtie} S = \sigma_C(R \times S)$

# Theta-Join

*S1*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

*R1*

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

$$S1 \bowtie_{S1.sid < R1.sid} R1$$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

$$R \bowtie_c S = \sigma_c (R \times S)$$

# Natural Join

- The natural join of two relations R and S is a set of **pairs of tuples**, one from R and one from S, **that agree on whatever attributes are common to the schemas of R and S**.
- The schema for the result contains the union of the attributes of R and S.
- Assume the schemas R(A,B, C) and S(B, C,D)

**RA:**  $R \bowtie S$

**SQL:** **SELECT \***

**FROM R , S**

**WHERE R.B = S.B AND R.C = S.C;**

# Operators Covered So far

- ▶ Remove parts of a single relation:
  - ▶ projection:  $\pi_{A,B}(R)$  and SELECT A, B FROM R.
  - ▶ selection:  $\sigma_C(R)$  and SELECT \* FROM R WHERE C.
  - ▶ combining projection and selection:
    - ▶  $\pi_{A,B}(\sigma_C(R))$
    - ▶ SELECT A, B FROM R WHERE C.
- ▶ Set operations ( $R$  and  $S$  must have the same attributes, same attribute types, and same order of attributes):
  - ▶ union:  $R \cup S$  and (R) UNION (S).
  - ▶ intersection:  $R \cap S$  and (R) INTERSECT (S).
  - ▶ difference:  $R - S$  and (R) EXCEPT (S).
- ▶ Combine the tuples of two relations:
  - ▶ Cartesian product:  $R \times S$  and ... FROM R, S ....
  - ▶ Theta-join:  $R \bowtie_C S$  and ... FROM R, S WHERE C.
  - ▶ Natural join:  $R \bowtie S$ ; in SQL, list the conditions that the common attributes be equal in the WHERE clause.



# Renaming

- If two relations have the same attribute, **disambiguate** the attributes **by prefixing the attribute with the name of the relation it belongs to.**
- How do we answer the query “Name pairs of students who live at the same address”? **Students(Name, Address)**
  - We need to take the cross-product of Students with itself?
  - How do we refer to the two “copies” of Students?
  - Use the rename operator.

**RA:**  $\rho_S(A_1, A_2, \dots, A_n)(R)$  : give R the name S; R has n attributes, which are called A1, A2, . . . , An in S

**SQL:** Use the **AS** keyword in the **FROM** clause: Students AS Students1  
renames Students to Students1.

**SQL:** Use the **AS** keyword in the **SELECT** clause to rename attributes.

# Renaming

Name pairs of students who live at the same address.

RA  $\pi_{S1.Name, S2.Name}(\sigma_{S1.Address = S2.Address}(\rho_{S1}(Students) \times \rho_{S2}(Students)))$ .

SQL `SELECT S1.name, S2.name  
FROM Students AS S1, Students AS S2  
WHERE S1.address = S2.address;`

- Are these correct ?
- **No !!!** the result includes tuples where a student is paired with himself/herself
- **Solution:** Add the condition `S1.name <> S2.name`.

# Practicing Relational Algebra

# Q1: Find names of sailors who have reserved boat #103

*Reserves*(sid, *bid*, day)

*Sailors*(sid, *sname*, rating, age)

- Solution 1:

$\pi_{sname}(\sigma_{bid = 103} (Reserves \bowtie Sailors))$

- Solution 2 (more efficient)

$\pi_{sname}((\sigma_{bid = 103} Reserves) \bowtie Sailors)$

- Solution 3 (using rename operator)

$P(Temp1 (\sigma_{bid = 103} Reserves))$

$P(Temp2 (Temp1 \bowtie Sailors))$

$\pi_{sname}(Temp2)$

## Q2: Find names of sailors who have reserved a red boat

*Reserves*(sid, *bid*, day)

*Sailors*(sid, *sname*, rating, age)

*Boats*(*bid*, bname, *color*)

- Solution 1:

$\pi_{sname}((\sigma_{color = 'red'} Boats) \bowtie Reserves \bowtie Sailors)$

- Solution 2 (more efficient)

$\pi_{sname}(\pi_{sid}((\pi_{bid} \sigma_{color = 'red'} Boats) \bowtie Reserves) \bowtie Sailors)$

## Q3: Find the colors of boats reserved by Lubber

*Reserves(sid, bid, day)*  
*Boats(bid, bname, color)*

*Sailors(sid, sname, rating, age)*

- Solution:

$\pi_{color}((\sigma_{sname = 'Lubber'} \mathbf{Sailor}) \bowtie \mathbf{Reserves} \bowtie \mathbf{Boats})$

Q4: Find the names of sailors who  
have reserved at least one boat

*Reserves(sid, bid, day)*  
*Boats(bid, bname, color)*

*Sailors(sid, **sname**, rating, age)*

- Solution:

$\pi_{sname}(\mathbf{Sailor} \bowtie \mathbf{Reserves})$

Q5: Find the names of sailors who have reserved a red or a green boat

*Reserves(sid, bid, day)*  
*Boats(bid, bname, color)*

*Sailors(sid, sname, rating, age)*

- Solution:

$\pi_{sname}(\sigma_{color='red' \text{ or } color = 'green'} \text{Boats} \bowtie \text{Reserves} \bowtie \text{Sailors})$



Q6: Find the names of sailors who have reserved a red and a green boat

*Reserves(sid, bid, day)*  
*Boats(bid, bname, color)*

*Sailors(sid, sname, rating, age)*

- Solution:

~~$\pi_{sname}(\sigma_{color='red' \text{ and } color='green'} \text{Boats} \bowtie \text{Reserves} \bowtie \text{Sailors})$~~

**A ship cannot have TWO colors at the same time**

$$\pi_{sname}(\sigma_{color='red'} \text{Boats} \bowtie \text{Reserves} \bowtie \text{Sailors})$$
$$\cap$$
$$\pi_{sname}(\sigma_{color='green'} \text{Boats} \bowtie \text{Reserves} \bowtie \text{Sailors})$$

Q7: Find the sids of sailors with **age over 20** who have not reserved a **red boat**

*Reserves*(sid, bid, day)  
*Boats*(bid, bname, **color**)

*Sailors*(**sid**, sname, rating, **age**)

Strategy ? ? ?

Find all sailors (sids) with age over 20

Find all sailors (sids) who have reserved a red boat

Take their set difference

- Solution:

$$\pi_{sid} (\sigma_{age>20} \text{ Sailors}) - \pi_{sid} ((\sigma_{color='red'} \text{ Boats}) \bowtie \text{ Reserves})$$